

Overfitting vs Underfitting

Overfitting

Overfitting occurs when a model learns the training data too well, memorizing noise and specific patterns that don't generalize to new data.

Characteristics:

- High performance on training data, poor performance on validation/test data
- Large gap between training and validation accuracy
- Model is too complex relative to the amount of training data
- Captures noise rather than underlying patterns

Signs of Overfitting:

- Training accuracy continues improving while validation accuracy plateaus or degrades
- Model performs exceptionally well on training set but poorly on new data
- Very complex decision boundaries that seem to "memorize" training examples

Common Causes:

- Too many parameters relative to training data size
- Training for too many epochs
- Insufficient regularization
- Very deep networks without proper constraints

Underfitting

Underfitting occurs when a model is too simple to capture the underlying patterns in the data.

Characteristics:

- Poor performance on both training and validation data
- Model lacks sufficient complexity to represent the data relationships
- High bias, low variance
- Overly simplistic assumptions about data structure

Signs of Underfitting:

- Both training and validation accuracy are low
- Model seems to ignore important features or relationships
- Performance doesn't improve significantly with more training

Common Causes:

- Model architecture too simple
- Insufficient training time
- Over-regularization
- Poor feature selection or engineering

The Bias-Variance Tradeoff

- **High Bias (Underfitting):** Model makes strong assumptions, misses relevant patterns
- **High Variance (Overfitting):** Model is sensitive to small changes in training data
- **Sweet Spot:** Balance between bias and variance for optimal generalization

Cross-Validation Basics

Cross-validation is a technique to assess how well a model generalizes to unseen data by systematically testing on different subsets of the available data.

K-Fold Cross-Validation

Process:

1. Split dataset into k equal-sized folds
2. For each fold: use $k-1$ folds for training, 1 fold for validation
3. Repeat k times, each time using a different fold for validation
4. Average the performance metrics across all k iterations

Benefits:

- More robust estimate of model performance
- Uses all data for both training and validation
- Reduces dependence on particular train/test split
- Helps detect overfitting

Common Variations:

- **5-fold or 10-fold:** Most common choices balancing computational cost and statistical reliability
- **Leave-One-Out (LOO):** k equals the number of samples; computationally expensive but maximizes training data
- **Stratified K-fold:** Maintains class distribution proportions in each fold

Time Series Cross-Validation

For temporal data where future cannot predict past:

- **Forward Chaining:** Train on data up to time t , test on $t+1$
- **Rolling Window:** Fixed-size training window that moves forward
- **Expanding Window:** Growing training set over time

Key Considerations

- **Data Leakage:** Ensure no information from validation set influences training
- **Computational Cost:** More folds = more training iterations
- **Statistical Significance:** More folds generally provide more reliable estimates

Model Improvement Techniques

Regularization Methods

L1 Regularization (Lasso):

- Adds penalty proportional to absolute value of parameters
- Encourages sparsity (some weights become exactly zero)
- Useful for feature selection
- Formula: $\text{Loss} + \lambda \sum |w_i|$

L2 Regularization (Ridge):

- Adds penalty proportional to square of parameters
- Shrinks weights toward zero but doesn't eliminate them
- Helps prevent overfitting by constraining parameter magnitude
- Formula: $\text{Loss} + \lambda \sum w_i^2$

Elastic Net:

- Combines L1 and L2 regularization
- Balances feature selection with weight shrinkage

Neural Network Specific Techniques

Dropout:

- Randomly sets a fraction of neurons to zero during training
- Forces network to not rely on specific neurons
- Different random subset dropped each training iteration
- Only applied during training, not inference

Batch Normalization:

- Normalizes inputs to each layer
- Reduces internal covariate shift
- Allows higher learning rates and faster training
- Acts as implicit regularization

Early Stopping:

- Monitor validation performance during training
- Stop training when validation performance stops improving
- Prevents overfitting from excessive training epochs
- Requires separate validation set

Data-Related Improvements

Data Augmentation:

- Artificially increase training data size
- Apply transformations that preserve labels
- Examples: rotation, scaling, cropping for images; synonym replacement for text

Feature Engineering:

- Create new features from existing ones

- Domain knowledge can guide feature creation
- Polynomial features, interaction terms, domain-specific transformations

Data Quality:

- Remove or handle outliers appropriately
- Address missing values systematically
- Ensure data consistency and proper preprocessing

Architecture and Training Improvements

Ensemble Methods:

- Combine predictions from multiple models
- Bagging: Train multiple models on different data subsets
- Boosting: Train models sequentially, focusing on previous errors
- Voting: Average or majority vote across model predictions

Hyperparameter Tuning:

- Grid Search: Exhaustive search over parameter combinations
- Random Search: Sample parameter combinations randomly
- Bayesian Optimization: Use probabilistic model to guide search
- Learning rate scheduling, architecture modifications

Transfer Learning:

- Start with pre-trained model from related domain
- Fine-tune on specific task with smaller dataset
- Particularly effective for computer vision and NLP tasks

Diagnostic Approach

1. **Identify the Problem:** Plot learning curves to distinguish overfitting vs underfitting
2. **Address Systematically:** Start with data quality, then model complexity, then regularization
3. **Validate Changes:** Use cross-validation to ensure improvements generalize
4. **Iterate:** Model improvement is typically an iterative process requiring multiple techniques